

EFFICIENT MAXIMUM FLOW ALGORITHM


Nikolay Sakharlykh - Hugo Braun; May 10, 2017



MAXIMUM FLOW

Definition

Example: How much instant power can Palo Alto get using that electric grid?





- Directed graph
- Flow capacities on edges
- Maximum flow from **s** to **t**?

MAXIMUM FLOW

Applications



Image segmentation



Community detection

MAXIMUM FLOW SOLVERS

AUGMENTING PATHS

Iteratively find a new augmenting path

- Ford-Fulkerson
- Edmonds-Karp
- Dinic's/MPM

PREFLOW

Push flow locally in a preflow graph

- Push relabel and its variants

OTHER

Linear programming

AGENDA


Edmonds-Karp

Push-relabel

MPM

FORD FULKERSON

Workflow




```
while a path p of capacity c > 0 exists from s to t:  
    maxflow += c  
    for all edges in path p:  
        edge.capacity -= c  
        add reverse edge from edge.destination to  
        edge.source of capacity c
```

→ Augmenting path on first iteration, path of capacity $\min(1, 3, 1) = 1$

FORD FULKERSON

Workflow




```
while a path p of capacity c > 0 exists from s to t:  
    maxflow += c  
    for all edges in path p:  
        edge.capacity -= c  
        add reverse edge from edge.destination to  
        edge.source of capacity c
```

→ Augmenting path on first iteration, path of capacity $\min(1, 3, 1) = 1$

FORD FULKERSON

Workflow




```
while a path p of capacity c > 0 exists from s to t:  
    maxflow += c  
    for all edges in path p:  
        edge.capacity -= c  
        add reverse edge from edge.destination to  
        edge.source of capacity c
```

—→ Augmenting path on first iteration, path of capacity $\min(1, 3, 1) = 1$

—→ Augmenting path on second iteration, path of capacity $\min(1, 1, 1, 3, 5) = 1$

EDMONDS-KARP



Edmonds-Karp: variation of Ford Fulkerson

Main idea: use *the shortest* augmenting path

One augmenting path needs one BFS

Wikipedia graph: ~5000 augmenting paths




Ford-Fulkerson & variants use too many graph traversals

PUSH-RELABEL

PUSH-RELABEL

Workflow



Saturate all out-edges of s , create reverse edges

$\ell[s] = \text{number of vertices}$

$\ell[v] = 0$ for all $v \in V \setminus \{s\}$

while there is an applicable push or relabel operation
execute the operation

push(u, v):

if($e[u] > 0$ and $\ell[u] == \ell[v] + 1$)

push $e[u]$ amount of flow from u to v


relabel(u):

if($e[u] > 0$ and $\ell[u] \leq \ell[v]$ for all current neighbors)

$\ell[u] = \text{minimum } \ell[v] \text{ among neighbors} + 1$

PUSH-RELABEL

Workflow



Saturate all out-edges of s , create reverse edges

$\ell[s] = \text{number of vertices}$

$\ell[v] = 0$ for all $v \in V \setminus \{s\}$

while there is an applicable push or relabel operation
execute the operation

push(u, v):

if($e[u] > 0$ and $\ell[u] == \ell[v] + 1$)

push $e[u]$ amount of flow from u to v


relabel(u):

if($e[u] > 0$ and $\ell[u] \leq \ell[v]$ for all current neighbors)

$\ell[u] = \text{minimum } \ell[v] \text{ among neighbors} + 1$

PUSH-RELABEL

Workflow



Saturate all out-edges of s , create reverse edges

$\ell[s] = \text{number of vertices}$

$\ell[v] = 0$ for all $v \in V \setminus \{s\}$

while there is an applicable push or relabel operation
execute the operation

push(u, v):

if($e[u] > 0$ and $\ell[u] == \ell[v] + 1$)

push $e[u]$ amount of flow from u to v


relabel(u):

if($e[u] > 0$ and $\ell[u] \leq \ell[v]$ for all current neighbors)

$\ell[u] = \text{minimum } \ell[v] \text{ among neighbors} + 1$

PUSH-RELABEL

Workflow



Saturate all out-edges of s , create reverse edges

$\ell[s] = \text{number of vertices}$

$\ell[v] = 0$ for all $v \in V \setminus \{s\}$

while there is an applicable push or relabel operation
execute the operation

push(u, v):

if($e[u] > 0$ and $\ell[u] == \ell[v] + 1$)

push $e[u]$ amount of flow from u to v

relabel(u):

if($e[u] > 0$ and $\ell[u] \leq \ell[v]$ for all current neighbors)

$\ell[u] = \text{minimum } \ell[v] \text{ among neighbors} + 1$

PUSH-RELABEL

Parallelism issues


Source of parallelism:

while there is an applicable push or relabel operation
execute the operation

At this step, we could relabel **a** or **d**. Which one?

Complexity of heuristics :

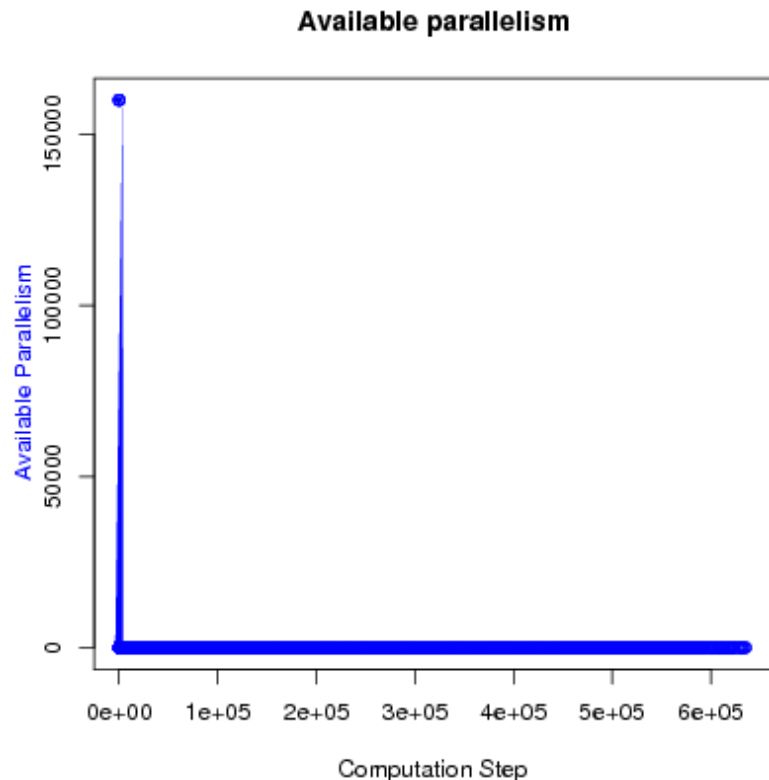
PRIORITY	LARGEST L	SMALLEST L	FIFO
Complexity	$O(V^2 \sqrt{E})$	$O(V^2 E)$	$O(V^3)$



Order affects convergence. Massive parallelism yields random order

PUSH-RELABEL

Parallelism issues



Source : The University of Texas at Austin

In theory,
number of threads = number of vertices


In practice,
number of active vertices \ll number of vertices



Parallelism drops. Not enough to saturate the GPU

PUSH-RELABEL

Conclusion



- Actual parallelism is **low**
- Massive parallelism yields **random order** which damages performance
- We need graph traversals (BFS) for some critical heuristics


road_usa: GPU does **20 BFS**, CPU does only **3 BFS**
CPU is faster since it requires fewer traversals

Push-relabel not suited for GPU implementation

MPM

DINIC'S

Workflow



Edges on paths of length 3

Two augmenting paths of length 3


They have been discovered using just one BFS

Avoid running BFS twice here

Main idea of Dinic's: *reuse* BFS results

DINIC'S

Workflow




Graph **G**

While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find augmenting path from s to t in G_L
Push corresponding flow in G_L , update edges

DINIC'S

Workflow




Graph $G_{L(3)}$

While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find augmenting path from s to t in G_L
Push corresponding flow in G_L , update edges

DINIC'S

Workflow



While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do


Find augmenting path from s to t in G_L

Push corresponding flow in G_L , update edges

DFS

DINIC'S

Workflow




Graph $G_{L(4)}$

While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find augmenting path from s to t in G_L
Push corresponding flow in G_L , update edges

DINIC'S

Workflow



Graph $G_{L(4)}$

```
While s is still connected to t in G, do
  Create layer graph  $G_L$  containing only
  shortest paths from s to t
```


```
While s is still connected to t, do
  Find augmenting path from s to t in  $G_L$ 
  Push corresponding flow in  $G_L$ , update edges
```

DFS

DFS traverse all vertices on GPU
We lose all advantages of Dinic's

MPM

Workflow



Graph **G**


While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find vertex u with minimum potential m , with
 $\text{potential}(u) = \min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))$

push m from u to t , pull m from s to u
remove all vertex with $\min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))=0$

MPM

Workflow



Graph **G**


While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find vertex u with minimum potential m , with
 $\text{potential}(u) = \min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))$

push m from u to t , pull m from s to u
remove all vertex with $\min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))=0$

MPM

Workflow



Graph $G_{L(3)}$

While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do

Find vertex u with minimum potential m , with
 $\text{potential}(u) = \min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))$


push m from u to t , pull m from s to u

remove all vertex with $\min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))=0$

c is selected so that we know 1 amount of flow will pass through

MPM

Workflow



Graph $G_{L(3)}$

While s is still connected to t in G , do
Create layer graph G_L containing only
shortest paths from s to t

While s is still connected to t , do
Find vertex u with minimum potential m , with
 $\text{potential}(u) = \min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))$


push m from u to t , pull m from s to u

remove all vertex with $\min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))=0$

c is selected so that we know 1 amount of flow will pass through

MPM

Workflow



Graph $G_L(3)$

```
While s is still connected to t in G, do
  Create layer graph  $G_L$  containing only
  shortest paths from s to t
```

```
While s is still connected to t, do
  Find vertex u with minimum potential m, with
   $\text{potential}(u) = \min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))$ 
```


```
  push m from u to t, pull m from s to u
```

```
  remove all vertex with  $\min(\text{degree}_{\text{in}}(u), \text{degree}_{\text{out}}(u))=0$ 
```


MPM

Dinic's vs MPM

Dinic's - DFS



MPM - Push/Pull/Prune




Across the graph, min potential = 1 (vertex **h**)
Pushing 1 to **t**, pulling 1 from **s**, using any edges

- Processed but useless edges
- Processed and acceptable edges

MPM

Dinic's vs MPM



Saturating one augmenting path on GPU:

- **MPM:** Push/pull/prune process **30us**
- **Edmonds-Karp/Dinic's:** one BFS **>1ms**


Perf. bounded by kernel launch latency

Example: Wikipedia 2011

- **MPM:** **5 BFS**, 6000 augmenting paths
- **EK:** **6000 BFS**

MPM

GPU design



MPM paper gives a high level implementation

Most of the work went into GPU
implementation design (2 out of 3 months)

MAXIMUM FLOW RESULTS

Galois on dual socket Haswell 16 cores vs NVIDIA Titan X (Pascal)

GRAPH	N	NNZ	SPEED UP		
			AVG	MIN	MAX
wiki03	455436	3811198	9.1	1.7	15.3
wiki11	3721339	121043107	22.5	19.8	28.9
road_usa	23947347	57708624	2	0.7	4.2
road CA	1971281	5533214	2.3	0.8	4.9

EFFICIENT MAXIMUM FLOW ALGORITHM

Takeaways

- **Black-box solver:** large variety of applications can be seen as the flow problem.
- **Data-dependent, irregular algorithm:** how to create enough “real” parallelism and how to avoid latency issues on the GPU.
- **Order of magnitude speed-ups on *wide graphs*.** Long graphs require a more efficient graph traversal implementation.

REFERENCES

- [An Experimental Comparison of Min-Cut/Max-Flow Algorithms for Energy Minimization in Vision](#), Yuri Boykov, Vladimir Kolmogorov
- [Finding Web Communities by Maximum Flow Algorithm using Well-Assigned Edge Capacities](#), Noriko IMAFUJI, and Masaru KITSUREGAWA
- [An \$O\(|V|^3\)\$ algorithm for finding maximum flows in networks](#), V.M. Malhotra, M.Pramodh Kumar, S.N. Maheshwari
- [Parallizing the Push-Relabel Max Flow Algorithm](#), Victoria Popic, Javier Vélez

